

Tidy data and data wrangling

Prof. Maria Tackett

Click for PDF of slides

Tidy data

Tidy data

Happy families are all alike; every unhappy family is unhappy in its own way.

Leo Tolstoy

Characteristics of tidy data:

- Each variable forms a column.
- Each observation forms a row.
- Each type of observational unit forms a table.

Characteristics of untidy data:

!@#\$%^&*()

What makes this data not tidy?

**Airplanes on Hand in the AAF, By Major Type:
Jul 1939 to Aug 1945**

End of Month	Total	Very Heavy Bombers	Heavy Bombers	Medium Bombers	Light Bombers	Fighters	Recon-naissance	Transports	Trainers	Communi-cations
1939										
Jul	2,402	-	16	400	276	494	356	118	735	7
Aug	2,440	-	18	414	276	492	359	129	745	7
[Germany invades Poland, 1 Sep 1939]										
Sep	2,473	-	22	428	278	489	359	136	754	7
Oct	2,507	-	27	446	277	490	365	137	758	7
Nov	2,536	-	32	458	275	498	375	136	755	7
Dec	2,546	-	39	464	274	492	378	131	761	7
1940										
Jan	2,588	-	45	466	271	464	409	128	798	7
Feb	2,658	-	49	470	271	458	415	128	860	7
Mar	2,709	-	54	468	267	453	415	125	920	7
Apr	2,806	-	54	468	263	451	416	125	1,022	7
May	2,906	-	54	470	259	459	410	124	1,123	7
Jun	2,966	-	54	478	166	477	414	127	1,243	7
[France surrenders to Germany, 25 Jun 1940]										
[Battle of Britain begins, 10 July 1940]										
Jul	3,102	-	56	483	161	500	410	128	1,357	7
Aug	3,295	-	65	485	158	539	407	128	1,506	7

What makes this data not tidy?

Subject	United States			
	Estimate	Margin of Error	Percent	Percent Margin of Error
EMPLOYMENT STATUS				
Population 16 years and over	255,797,692	+/-17,051	255,797,692	(X)
In labor force	162,184,325	+/-135,158	63.4%	+/-0.1
Civilian labor force	161,159,470	+/-127,501	63.0%	+/-0.1
Employed	150,599,165	+/-138,066	58.9%	+/-0.1
Unemployed	10,560,305	+/-27,385	4.1%	+/-0.1
Armed Forces	1,024,855	+/-10,363	0.4%	+/-0.1
Not in labor force	93,613,367	+/-126,007	36.6%	+/-0.1
Civilian labor force	161,159,470	+/-127,501	161,159,470	(X)
Unemployment Rate	(X)	(X)	6.6%	+/-0.1
Females 16 years and over	131,092,196	+/-11,187	131,092,196	(X)
In labor force	76,493,327	+/-75,824	58.4%	+/-0.1
Civilian labor force	76,350,498	+/-75,238	58.2%	+/-0.1
Employed	71,451,559	+/-79,007	54.5%	+/-0.1
Own children of the householder under 6 years	22,939,897	+/-14,240	22,939,897	(X)
All parents in family in labor force	14,957,537	+/-36,506	65.2%	+/-0.1
Own children of the householder 6 to 17 years	47,007,147	+/-19,644	47,007,147	(X)
All parents in family in labor force	33,238,793	+/-49,036	70.7%	+/-0.1

Summary tables

Is each of the following a dataset or a summary table?

```
## # A tibble: 87 x 3
```

##	name	height	mass
##	<chr>	<int>	<dbl>
##	1 Luke Skywalker	172	77
##	2 C-3PO	167	75
##	3 R2-D2	96	32
##	4 Darth Vader	202	136
##	5 Leia Organa	150	49
##	6 Owen Lars	178	120
##	7 Beru Whitesun lars	165	75
##	8 R5-D4	97	32
##	9 Biggs Darklighter	183	84
##	10 Obi-Wan Kenobi	182	77
##	... with 77 more rows		

```
## # A tibble: 3 x 2
```

##	gender	avg_height
##	<chr>	<dbl>
##	1 feminine	165.
##	2 masculine	177.
##	3 <NA>	181.

Displaying data

```
starwars %>%  
  select(name, height, mass)
```


Summarizing data

```
starwars %>%  
  group_by(gender) %>%  
  summarize(  
    avg_height = mean(height, na.rm = TRUE) %>% round(2)  
  )
```


Grammar of data wrangling

A grammar of data wrangling...

... based on the concepts of functions as verbs that manipulate data frames



- **select**: pick columns by name
- **arrange**: reorder rows
- **slice**: pick rows using index(es)
- **filter**: pick rows matching criteria
- **distinct**: filter for unique rows
- **mutate**: add new variables
- **summarise**: reduce variables to values
- **group_by**: for grouped operations

Rules of dplyr functions

- First argument is *always* a data frame
- Subsequent arguments say what to do with that data frame
- Always return a data frame

Data: Hotel bookings

- Data from two hotels: one resort and one city hotel
- Observations: Each row represents a hotel booking
- Goal for original data collection: Development of prediction models to classify a hotel booking's likelihood to be cancelled ([Antonia et al., 2019](#))
- Featured in [TidyTuesday](#)!

```
hotels <- read_csv("data/hotels.csv")
```

First look: Variables

```
names(hotels)
```

```
## [1] "hotel" "is_canceled"
## [3] "lead_time" "arrival_date_year"
## [5] "arrival_date_month" "arrival_date_week_number"
## [7] "arrival_date_day_of_month" "stays_in_weekend_nights"
## [9] "stays_in_week_nights" "adults"
## [11] "children" "babies"
## [13] "meal" "country"
## [15] "market_segment" "distribution_channel"
## [17] "is_repeated_guest" "previous_cancellations"
## [19] "previous_bookings_not_canceled" "reserved_room_type"
## [21] "assigned_room_type" "booking_changes"
## [23] "deposit_type" "agent"
## [25] "company" "days_in_waiting_list"
## [27] "customer_type" "adr"
```

Second look: Overview

```
glimpse(hotels)
```

```
## Rows: 119,390
```

```
## Columns: 32
```

```
## $ hotel
```

```
## $ is_canceled
```

```
## $ lead_time
```

```
## $ arrival_date_year
```

```
## $ arrival_date_month
```

```
## $ arrival_date_week_number
```

```
## $ arrival_date_day_of_month
```

```
## $ stays_in_weekend_nights
```

```
## $ stays_in_week_nights
```

```
## $ adults
```

```
## $ children
```

```
## $ babies
```

```
<chr> "Resort Hotel", "Resort Hotel", "Res
```

```
<dbl> 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0,
```

```
<dbl> 342, 737, 7, 13, 14, 14, 0, 9, 85, 7
```

```
<dbl> 2015, 2015, 2015, 2015, 2015, 2015,
```

```
<chr> "July", "July", "July", "July", "Ju"
```

```
<dbl> 27, 27, 27, 27, 27, 27, 27, 27, 27,
```

```
<dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
```

```
<dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
```

```
<dbl> 0, 0, 1, 1, 2, 2, 2, 2, 3, 3, 4, 4,
```

```
<dbl> 2, 2, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2,
```

```
<dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
```

```
<dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
```



Select a single column

View only the **lead_time** (number of days between booking and arrival date):

```
hotels %>%  
  select(lead_time)
```

```
## # A tibble: 119,390 x 1  
##   lead_time  
##   <dbl>  
## 1      342  
## 2      737  
## 3         7  
## 4        13  
## 5        14  
## 6        14  
## 7         0  
## 8         0
```

- Start with a data frame
- Pass it to the **select()** function.
- Second argument is variable we want to select: **lead_time**
- The result is a data frame with 119,300 and 1 column: --dplyr functions always expect a data frame and always yield a data frame.

Select multiple columns

View only the **hotel** type and **lead_time**:

```
hotels %>%  
  select(hotel, lead_time)
```

```
## # A tibble: 119,390 x 2  
##   hotel      lead_time  
##   <chr>      <dbl>  
## 1 Resort Hotel      342  
## 2 Resort Hotel      737  
## 3 Resort Hotel        7  
## 4 Resort Hotel      13  
## 5 Resort Hotel      14  
## 6 Resort Hotel      14  
## 7 Resort Hotel        0  
## 8 Resort Hotel        0
```

What if we wanted to select these columns, and then arrange the data in descending order of lead time?

Data wrangling, step-by-step

Select:

```
hotels %>%  
  select(hotel, lead_time)
```

```
## # A tibble: 119,390 x 2  
##   hotel      lead_time  
##   <chr>      <dbl>  
## 1 Resort Hotel      342  
## 2 Resort Hotel      737  
## 3 Resort Hotel        7  
## 4 Resort Hotel      13  
## 5 Resort Hotel      14  
## 6 Resort Hotel      14  
## 7 Resort Hotel        0  
## 8 Resort Hotel        0
```

Select, then arrange:

```
hotels %>%  
  select(hotel, lead_time) %>%  
  arrange(desc(lead_time))
```

```
## # A tibble: 119,390 x 2  
##   hotel      lead_time  
##   <chr>      <dbl>  
## 1 Resort Hotel      737  
## 2 Resort Hotel      709  
## 3 City Hotel        629  
## 4 City Hotel        629  
## 5 City Hotel        629  
## 6 City Hotel        629  
## 7 City Hotel        629
```

Pipes

What is a pipe?

In programming, a pipe is a technique for passing information from one process to another.

- Start with the data frame **hotels**, and pass it to the **select()** function,

```
hotels %>%  
  select(hotel, lead_time) %>%  
  arrange(desc(lead_time))
```

```
## # A tibble: 119,390 x 2  
##   hotel      lead_time  
##   <chr>      <dbl>  
## 1 Resort Hotel      737  
## 2 Resort Hotel      709  
## 3 City Hotel        629  
## 4 City Hotel        629  
## 5 City Hotel        629  
## 6 City Hotel        629  
## 7 City Hotel        629  
## 8 City Hotel        629  
## 9 City Hotel        629  
## 10 City Hotel       629  
## # ... with 119,380 more rows
```

What is a pipe?

In programming, a pipe is a technique for passing information from one process to another.

- Start with the data frame **hotels**, and pass it to the **select()** function,
- then we select the variables **hotel** and **lead_time**,

```
hotels %>%  
  select(hotel, lead_time) %>%  
  arrange(desc(lead_time))
```

```
## # A tibble: 119,390 x 2  
##   hotel      lead_time  
##   <chr>      <dbl>  
## 1 Resort Hotel      737  
## 2 Resort Hotel      709  
## 3 City Hotel        629  
## 4 City Hotel        629  
## 5 City Hotel        629  
## 6 City Hotel        629  
## 7 City Hotel        629  
## 8 City Hotel        629  
## 9 City Hotel        629  
## 10 City Hotel        629  
## # ... with 119,380 more rows
```

What is a pipe?

In programming, a pipe is a technique for passing information from one process to another.

- Start with the data frame **hotels**, and pass it to the **select()** function,
- then we select the variables **hotel** and **lead_time**,
- and then we arrange the data frame by **lead_time** in descending order.

```
hotels %>%  
  select(hotel, lead_time) %>%  
  arrange(desc(lead_time))
```

```
## # A tibble: 119,390 x 2  
##   hotel      lead_time  
##   <chr>      <dbl>  
## 1 Resort Hotel      737  
## 2 Resort Hotel      709  
## 3 City Hotel        629  
## 4 City Hotel        629  
## 5 City Hotel        629  
## 6 City Hotel        629  
## 7 City Hotel        629  
## 8 City Hotel        629  
## 9 City Hotel        629  
## 10 City Hotel        629  
## # ... with 119,380 more rows
```

Aside

The pipe operator is implemented in the package **magrittr**, though we don't need to load this package explicitly since **tidyverse** does this for us.

Any guesses as to why the package is called magrittr?



How does a pipe work?

- You can think about the following sequence of actions - find keys, start car, drive to work, park.
- Expressed as a set of nested functions in R pseudocode this would look like:

```
park(drive(start_car(find("keys")), to = "work"))
```

- Writing it out using pipes give it a more natural (and easier to read) structure:
 - Read the pipe as "and then"

```
find("keys") %>%  
  start_car() %>%  
  drive(to = "work") %>%  
  park()
```

What about other arguments?

Use the dot to

- send results to a function argument other than first one or
- use the previous result for multiple arguments

```
hotels %>%  
  filter(hotel == "Resort Hotel") %>%  
  lm(adr ~ lead_time, data = .)
```

```
##  
## Call:  
## lm(formula = adr ~ lead_time, data = .)  
##  
## Coefficients:  
## (Intercept)      lead_time  
##      93.16876       0.01925
```


Working with a single data frame

You have a single data frame, and you want to process it and prepare it for analysis!

select to keep variables

```
hotels %>%
```

```
  select(hotel, lead_time)
```

```
## # A tibble: 119,390 x 2
##   hotel      lead_time
##   <chr>      <dbl>
## 1 Resort Hotel      342
## 2 Resort Hotel      737
## 3 Resort Hotel        7
## 4 Resort Hotel      13
## 5 Resort Hotel      14
## 6 Resort Hotel      14
## 7 Resort Hotel        0
## 8 Resort Hotel        9
## 9 Resort Hotel      85
## 10 Resort Hotel      75
```

select to exclude variables

```
hotels %>%  
  select(-agent)
```

```
## # A tibble: 119,390 x 31  
##   hotel is_canceled lead_time arrival_date_ye... arrival_date_mo...  
##   <chr>      <dbl>      <dbl>          <dbl> <chr>  
## 1 Reso...      0      342          2015 July  
## 2 Reso...      0      737          2015 July  
## 3 Reso...      0       7          2015 July  
## 4 Reso...      0      13          2015 July  
## 5 Reso...      0      14          2015 July  
## 6 Reso...      0      14          2015 July  
## 7 Reso...      0       0          2015 July  
## 8 Reso...      0       9          2015 July  
## 9 Reso...      1      85          2015 July  
## 10 Reso...     1      75          2015 July  
## # ... with 119,380 more rows, and 26 more variables:  
## #   arrival_date_week_number <dbl>, arrival_date_day_of_month <dbl>,  
## #   stays_in_weekend_nights <dbl>, stays_in_week_nights <dbl>, adults <dbl>,  
## #   children <dbl>, babies <dbl>, meal <chr>, country <chr>,  
## #   market_segment <chr>, distribution_channel <chr>, is_repeated_guest <dbl>,  
## #   previous_cancellations <dbl>, previous_bookings_not_canceled <dbl>,  
## #   reserved_room_type <chr>, assigned_room_type <chr>, booking_changes <dbl>,  
## #   ...
```

select a range of variables

```
hotels %>%
```

```
  select(hotel:arrival_date_month)
```

```
## # A tibble: 119,390 x 5
```

```
##   hotel          is_canceled lead_time arrival_date_year arrival_date_month
##   <chr>          <dbl>      <dbl>          <dbl> <chr>
## 1 Resort Hotel      0        342          2015 July
## 2 Resort Hotel      0        737          2015 July
## 3 Resort Hotel      0         7          2015 July
## 4 Resort Hotel      0        13          2015 July
## 5 Resort Hotel      0        14          2015 July
## 6 Resort Hotel      0        14          2015 July
## 7 Resort Hotel      0         0          2015 July
## 8 Resort Hotel      0         9          2015 July
## 9 Resort Hotel      1        85          2015 July
## 10 Resort Hotel     1        75          2015 July
```

arrange in ascending / descending order

```
hotels %>%  
  select(adults, children, babies)  
  arrange(babies)
```

```
## # A tibble: 119,390 x 3  
##   adults children babies  
##   <dbl>     <dbl> <dbl>  
## 1         2         0     0  
## 2         2         0     0  
## 3         1         0     0  
## 4         1         0     0  
## 5         2         0     0  
## 6         2         0     0  
## 7         2         0     0  
## 8         2         0     0  
## 9         2         0     0
```

```
hotels %>%  
  select(adults, children, babies)  
  arrange(desc(babies))
```

```
## # A tibble: 119,390 x 3  
##   adults children babies  
##   <dbl>     <dbl> <dbl>  
## 1         2         0    10  
## 2         1         0     9  
## 3         2         0     2  
## 4         2         0     2  
## 5         2         0     2  
## 6         2         0     2  
## 7         2         0     2  
## 8         2         0     2  
## 9         2         0     2
```

slice for certain row numbers

```
# first five
```

```
hotels %>%
```

```
  slice(1:5)
```

```
## # A tibble: 5 x 32
```

```
##   hotel is_canceled lead_time arrival_date_ye... arrival_date_mo... arrival_date_we...
```

```
##   <chr>         <dbl>    <dbl>          <dbl> <chr>                <dbl>
```

```
## 1 Reso...         0      342          2015 July                27
```

```
## 2 Reso...         0      737          2015 July                27
```

```
## 3 Reso...         0        7          2015 July                27
```

```
## 4 Reso...         0       13          2015 July                27
```

```
## 5 Reso...         0       14          2015 July                27
```

```
## # ... with 26 more variables: arrival_date_day_of_month <dbl>,
```

```
## #   stays_in_weekend_nights <dbl>, stays_in_week_nights <dbl>, adults <dbl>,
```

```
## #   children <dbl>, babies <dbl>, meal <chr>, country <chr>,
```

```
## #   market_segment <chr>, distribution_channel <chr>, is_repeated_guest <dbl>,
```

```
## #   previous_cancellations <dbl>, previous_bookings_not_canceled <dbl>,
```

```
## #   reserved_room_type <chr>, assigned_room_type <chr>, booking_changes <dbl>,
```

```
## #   deposit_type <chr>, agent <chr>, company <chr>, days_in_waiting_list <dbl>,
```

```
## #   customer_type <chr>, adr <dbl>, required_car_parking_spaces <dbl>
```

Tip:

In R, you can use the `#` (hashtag or pound sign, depending on your age 🤪) for adding comments to your code. Any text following `#` will be printed as is, and won't be run as R code. This is useful for leaving comments in your code and for temporarily disabling certain lines of code while debugging.

```
hotels %>%  
  # slice the first five rows # this line is a comment  
  #select(hotel) %>%          # this one doesn't run  
  slice(1:5)                  # this line runs
```

```
## # A tibble: 5 x 32  
##   hotel is_canceled lead_time arrival_date_ye... arrival_date_mo... arrival_date_we...  
##   <chr>      <dbl>      <dbl>          <dbl> <chr>              <dbl>  
## 1 Reso...      0        342          2015 July              27  
## 2 Reso...      0        737          2015 July              27  
## 3 Reso...      0         7          2015 July              27  
## 4 Reso...      0        13          2015 July              27  
## 5 Reso...      0        14          2015 July              27  
## # ... with 26 more variables: arrival_date_day_of_month <dbl>,  
## #   stays_in_weekend_nights <dbl>, stays_in_week_nights <dbl>, adults <dbl>,
```


slice for certain row numbers

```
# last five
last_row <- nrow(hotels) # nrow() gives the number of rows in a data frame
hotels %>%
  slice((last_row - 4):last_row)
```

```
## # A tibble: 5 x 32
##   hotel is_canceled lead_time arrival_date_ye... arrival_date_mo... arrival_date_we...
##   <chr>      <dbl>      <dbl>          <dbl> <chr>              <dbl>
## 1 City...      0        23          2017 August              35
## 2 City...      0       102          2017 August              35
## 3 City...      0        34          2017 August              35
## 4 City...      0       109          2017 August              35
## 5 City...      0       205          2017 August              35
## # ... with 26 more variables: arrival_date_day_of_month <dbl>,
## #   stays_in_weekend_nights <dbl>, stays_in_week_nights <dbl>, adults <dbl>,
## #   children <dbl>, babies <dbl>, meal <chr>, country <chr>,
## #   market_segment <chr>, distribution_channel <chr>, is_repeated_guest <dbl>,
## #   previous_cancellations <dbl>, previous_bookings_not_canceled <dbl>,
## #   reserved_room_type <chr>, assigned_room_type <chr>, booking_changes <dbl>,
## #   deposit_type <chr>, agent <chr>, company <chr>, days_in_waiting_list <dbl>
```

filter to select a subset of rows

```
# bookings in City Hotels
```

```
hotels %>%
```

```
  filter(hotel == "City Hotel")
```

```
## # A tibble: 79,330 x 32
```

```
##   hotel is_canceled lead_time arrival_date_ye... arrival_date_mo...
```

```
##   <chr>         <dbl>      <dbl>          <dbl> <chr>
```

```
##  1 City...           0          6          2015 July
```

```
##  2 City...           1         88          2015 July
```

```
##  3 City...           1         65          2015 July
```

```
##  4 City...           1         92          2015 July
```

```
##  5 City...           1        100          2015 July
```

```
##  6 City...           1         79          2015 July
```

```
##  7 City...           0          3          2015 July
```

```
##  8 City...           1         63          2015 July
```

```
##  9 City...           1         62          2015 July
```

```
## 10 City...           1         62          2015 July
```

```
## # ... with 79,320 more rows, and 27 more variables:
```

```
## #   arrival_date_week_number <dbl>, arrival_date_day_of_month <dbl>,
```

```
## #   stays_in_weekend_nights <dbl> stays_in_week_nights <dbl> adults <dbl>
```

filter for many conditions at once

```
hotels %>%  
  filter(  
    adults == 0,  
    children >= 1  
  ) %>%  
  select(adults, babies, children)
```

```
## # A tibble: 223 x 3  
##   adults babies children  
##   <dbl>   <dbl>   <dbl>  
## 1      0      0       3  
## 2      0      0       2  
## 3      0      0       2  
## 4      0      0       2  
## 5      0      0       2  
## 6      0      0       3
```

filter for more complex conditions

```
# bookings with no adults and some children or babies in the room
hotels %>%
  filter(
    adults == 0,
    children >= 1 | babies >= 1      # / means or
  ) %>%
  select(adults, babies, children)
```

```
## # A tibble: 223 x 3
##   adults babies children
##   <dbl>   <dbl>   <dbl>
## 1      0      0       3
## 2      0      0       2
## 3      0      0       2
## 4      0      0       2
## 5      0      0       2
```

Logical operators in R

operator	definition	operator	definition
<	less than	x y	x OR y
<=	less than or equal to	is.na(x)	test if x is NA
>	greater than	!is.na(x)	test if x is not NA
>=	greater than or equal to	x %in% y	test if x is in y
==	exactly equal to	!(x %in% y)	test if x is not in y
!=	not equal to	!x	not x
x & y	x AND y		

Demo

distinct to filter for unique rows

... and **arrange** to order alphabetically

```
hotels %>%  
  distinct(market_segment) %>%  
  arrange(market_segment)
```

```
## # A tibble: 8 x 1  
##   market_segment  
##   <chr>  
## 1 Aviation  
## 2 Complementary  
## 3 Corporate  
## 4 Direct  
## 5 Groups  
## 6 Offline TA/TO  
## 7 Online TA  
## 8 Undefined
```

```
hotels %>%  
  distinct(hotel, market_segment) %>%  
  arrange(hotel, market_segment)
```

```
## # A tibble: 14 x 2  
##   hotel          market_segment  
##   <chr>         <chr>  
## 1 City Hotel    Aviation  
## 2 City Hotel    Complementary  
## 3 City Hotel    Corporate  
## 4 City Hotel    Direct  
## 5 City Hotel    Groups  
## 6 City Hotel    Offline TA/TO  
## 7 City Hotel    Online TA  
## 8 City Hotel    Undefined  
## 9 Resort Hotel  Complementary  
## 10 Resort Hotel Corporate  
## 11 Resort Hotel Direct  
## 12 Resort Hotel Groups  
## 13 Resort Hotel Offline TA/TO
```

count to create frequency tables

```
# alphabetical order by default
hotels %>%
  count(market_segment)
```

```
## # A tibble: 8 x 2
##   market_segment      n
##   <chr>          <int>
## 1 Aviation         237
## 2 Complementary     743
## 3 Corporate        5295
## 4 Direct          12606
## 5 Groups           19811
## 6 Offline TA/T0    24219
## 7 Online TA        56477
## 8 Undefined         2
```

```
# descending frequency order
hotels %>%
  count(market_segment,
        sort = TRUE)
```

```
## # A tibble: 8 x 2
##   market_segment      n
##   <chr>          <int>
## 1 Online TA        56477
## 2 Offline TA/T0    24219
## 3 Groups           19811
## 4 Direct          12606
## 5 Corporate        5295
## 6 Complementary     743
## 7 Aviation         237
## 8 Undefined         2
```


count and arrange

```
# ascending frequency order
hotels %>%
  count(market_segment) %>%
  arrange(n)
```

```
## # A tibble: 8 x 2
##   market_segment      n
##   <chr>          <int>
## 1 Undefined           2
## 2 Aviation          237
## 3 Complementary      743
## 4 Corporate          5295
## 5 Direct            12606
## 6 Groups            19811
## 7 Offline TA/T0      24219
## 8 Online TA          56477
```

```
# descending frequency order
# just like adding sort = TRUE
hotels %>%
  count(market_segment) %>%
  arrange(desc(n))
```

```
## # A tibble: 8 x 2
##   market_segment      n
##   <chr>          <int>
## 1 Online TA          56477
## 2 Offline TA/T0      24219
## 3 Groups            19811
## 4 Direct            12606
## 5 Corporate          5295
## 6 Complementary      743
## 7 Aviation           237
```

count for multiple variables

```
hotels %>%
```

```
  count(hotel, market_segment)
```

```
## # A tibble: 14 x 3
```

##	hotel	market_segment	n
##	<chr>	<chr>	<int>
## 1	City Hotel	Aviation	237
## 2	City Hotel	Complementary	542
## 3	City Hotel	Corporate	2986
## 4	City Hotel	Direct	6093
## 5	City Hotel	Groups	13975
## 6	City Hotel	Offline TA/TO	16747
## 7	City Hotel	Online TA	38748
## 8	City Hotel	Undefined	2
## 9	Resort Hotel	Complementary	201
## 10	Resort Hotel	Corporate	2309

order matters when you count

```
# hotel type first
```

```
hotels %>%
```

```
  count(hotel, market_segment)
```

```
## # A tibble: 14 x 3
```

```
##   hotel      market_segment      n
```

```
##   <chr>      <chr>          <int>
```

```
## 1 City Hotel Aviation           237
```

```
## 2 City Hotel Complementary       542
```

```
## 3 City Hotel Corporate          2986
```

```
## 4 City Hotel Direct             6093
```

```
## 5 City Hotel Groups            13975
```

```
## 6 City Hotel Offline TA/TO      16747
```

```
## 7 City Hotel Online TA          38748
```

```
## 8 City Hotel Undefined           2
```

```
## 9 Resort Hotel Complementary      201
```

```
## 10 Resort Hotel Corporate         2309
```

```
## 11 Resort Hotel Direct            6513
```

```
## 12 Resort Hotel Groups           5836
```

```
# market segment first
```

```
hotels %>%
```

```
  count(market_segment, hotel)
```

```
## # A tibble: 14 x 3
```

```
##   market_segment hotel      n
```

```
##   <chr>      <chr>          <int>
```

```
## 1 Aviation    City Hotel      237
```

```
## 2 Complementary City Hotel      542
```

```
## 3 Complementary Resort Hotel    201
```

```
## 4 Corporate   City Hotel      2986
```

```
## 5 Corporate   Resort Hotel    2309
```

```
## 6 Direct      City Hotel      6093
```

```
## 7 Direct      Resort Hotel    6513
```

```
## 8 Groups      City Hotel     13975
```

```
## 9 Groups      Resort Hotel    5836
```

```
## 10 Offline TA/TO City Hotel    16747
```

```
## 11 Offline TA/TO Resort Hotel    7472
```

```
## 12 Online TA   City Hotel     38748
```

Demo

mutate to add a new variable

```
hotels %>%  
  mutate(little_ones = children + babies) %>%  
  select(children, babies, little_ones) %>%  
  arrange(desc(little_ones))
```

```
## # A tibble: 119,390 x 3  
##   children babies little_ones  
##   <dbl>   <dbl>     <dbl>  
## 1      10      0         10  
## 2       0     10         10  
## 3       0      9          9  
## 4       2      1          3  
## 5       2      1          3  
## 6       2      1          3  
## 7       3      0          3  
## 8       2      1          3
```

Little ones in resort and city hotels

```
# Resort Hotel
hotels %>%
  mutate(little_ones = children + babies) %>%
  filter(
    little_ones >= 1,
    hotel == "Resort Hotel"
  ) %>%
  select(hotel, little_ones)
```

```
## # A tibble: 3,929 x 2
##   hotel      little_ones
##   <chr>         <dbl>
## 1 Resort Hotel         1
## 2 Resort Hotel         2
## 3 Resort Hotel         2
## 4 Resort Hotel         2
## 5 Resort Hotel         1
## 6 Resort Hotel         1
## 7 Resort Hotel         2
## 8 Resort Hotel         2
## 9 Resort Hotel         1
## 10 Resort Hotel        1
## #   with 3 919 more rows
```

```
# City Hotel
hotels %>%
  mutate(little_ones = children + babies) %>%
  filter(
    little_ones > 1,
    hotel == "City Hotel"
  ) %>%
  select(hotel, little_ones)
```

```
## # A tibble: 2,140 x 2
##   hotel      little_ones
##   <chr>         <dbl>
## 1 City Hotel         2
## 2 City Hotel         2
## 3 City Hotel         2
## 4 City Hotel         2
## 5 City Hotel         2
## 6 City Hotel         2
## 7 City Hotel         2
## 8 City Hotel         2
## 9 City Hotel         2
## 10 City Hotel        3
## #   with 2 130 more rows
```

What is happening in the following chunk?

```
hotels %>%  
  mutate(little_ones = children + babies) %>%  
  count(hotel, little_ones) %>%  
  mutate(prop = n / sum(n))
```

```
## # A tibble: 12 x 4  
##   hotel      little_ones      n      prop  
##   <chr>      <dbl> <int>    <dbl>  
## 1 City Hotel      0 73923 0.619  
## 2 City Hotel      1  3263 0.0273  
## 3 City Hotel      2  2056 0.0172  
## 4 City Hotel      3    82 0.000687  
## 5 City Hotel      9     1 0.00000838  
## 6 City Hotel     10     1 0.00000838  
## 7 City Hotel     NA     4 0.0000335  
## 8 Resort Hotel    0 36131 0.303  
## 9 Resort Hotel    1  2183 0.0183  
## 10 Resort Hotel   2  1716 0.0144
```

summarise for summary stats

```
# mean average daily rate for all bookings
```

```
hotels %>%
```

```
  summarise(mean_adr = mean(adr))
```

```
## # A tibble: 1 x 1
```

```
##   mean_adr
```

```
##      <dbl>
```

```
## 1      102.
```

Tip:

summarise() changes the data frame entirely, it collapses rows down to a single summary statistics, and removes all columns that are irrelevant to the calculation.

Tip:

summarise() also lets you get away with being sloppy and not naming your new column, but that's not recommended!



```
hotels %>%  
  summarise(mean(adr))
```

```
## # A tibble: 1 x 1  
##   `mean(adr)`  
##   <dbl>  
## 1      102.
```



```
hotels %>%  
  summarise(mean_adr = mean(adr))
```

```
## # A tibble: 1 x 1  
##   mean_adr  
##   <dbl>  
## 1      102.
```

group_by for grouped operations

```
# mean average daily rate for all booking at city and resort hotels  
hotels %>%
```

```
  group_by(hotel) %>%  
    summarise(mean_adr = mean(adr))
```

```
## # A tibble: 2 x 2  
##   hotel      mean_adr  
##   <chr>      <dbl>  
## 1 City Hotel    105.  
## 2 Resort Hotel   95.0
```

Calculating frequencies

The following two give the same result, so **count** is simply short for **group_by** then determine frequencies

```
hotels %>%  
  group_by(hotel) %>%  
  summarise(n = n())
```

```
## # A tibble: 2 x 2  
##   hotel          n  
##   <chr>      <int>  
## 1 City Hotel  79330  
## 2 Resort Hotel 40060
```

```
hotels %>%  
  count(hotel)
```

```
## # A tibble: 2 x 2  
##   hotel          n  
##   <chr>      <int>  
## 1 City Hotel  79330  
## 2 Resort Hotel 40060
```

Multiple summary statistics

summarise can be used for multiple summary statistics as well

```
hotels %>%  
  summarise(  
    min_adr = min(adr),  
    mean_adr = mean(adr),  
    median_adr = median(adr),  
    max_adr = max(adr)  
  )
```

```
## # A tibble: 1 x 4  
##   min_adr mean_adr median_adr max_adr  
##   <dbl>   <dbl>   <dbl>   <dbl>  
## 1    -6.38    102.    94.6    5400
```

Demo